

Revisiting the Semi-parametric Latent Factor Model with Regression Trees

Sameer K. Deshpande

March 6, 2019

1 General Setup

As a motivating example, consider modeling q physiological time series, which may be highly interdependent and which may be irregularly sampled (i.e. at any one time we may not observe realizations from each series). The main goal is to impute the value of each time series, to forecast each series several steps into the future, and to provide honest uncertainty quantification about these projections.

Formally, suppose we observe triplets of data $(\mathbf{x}_1, \mathbf{y}_1, \delta_1), \dots, (\mathbf{x}_n, \mathbf{y}_n, \delta_n)$ where $\mathbf{x}_i \in \mathbb{R}^p$ are covariates (possibly time-dependent), $\mathbf{y}_i \in \mathbb{R}^q$ are the noisy observations from each series, and $\delta_i \in \{0, 1\}^q$ is a vector of indicators with $\delta_{i,k} = 1$ if and only if $y_{i,k}$ is observed. For now, we assume that δ is deterministic. [\[skd\]: in the context of medical time series, this assumption is not particularly realistic. To wit, one series may track the results of a particular diagnostic test that is run only if other vital signs are in some critical region. That is, it's not unrealistic to believe that we observe one series based on the values of previous series. We'll return to this issue later](#)

Independent of δ_i , we model for each $i = 1, \dots, n$ and $k = 1, \dots, q$

$$y_{i,k} = f_k(\mathbf{x}_i) + \sigma_k(\mathbf{x}_i)\varepsilon_{i,k}$$

where $\mathbf{f} = (f_1, \dots, f_q)$ and $\boldsymbol{\sigma} = (\sigma_1, \dots, \sigma_q)$ are unknown vector valued functions of \mathbf{x} and the $\varepsilon_{i,k}$'s are independent standard normals. For now, we focus on the homoskedastic case

where the residual variances are constant and do not depend on inputs \mathbf{x}_i ; we will return to the heteroskedastic case later.

We take a Bayesian approach, which amounts to specifying a prior over the functions \mathbf{f} and σ and updating them with Bayes’ theorem to get posterior distributions that reflect our uncertainty about them in light of the data. In principle, we can attempt to learn each f_k independently of one another in an embarrassingly parallel manner. This, of course, precludes “sharing of statistical strength” and we can quite reasonably expect better predictive performance if we take advantage of the potential correlations between the outcomes (see, e.g., [Breiman and Friedman, 1997](#)). Improving prediction of one outcome dimension using information from some or all of the other outcome dimensions has been well-studied in machine learning community, under the name of “transfer learning” or “multi-task learning.” We borrow from this literature, and refer to each outcome dimensions as a task below.

This is a growing literature that use multi-task GPs regression to model multiple physiological time series (see, e.g., [Clifton et al. \(2012\)](#), [Ghassemi et al. \(2015\)](#), [Durichen et al. \(2015\)](#), [Futoma et al. \(2017\)](#), [Colopy et al. \(2018\)](#), and [Cheng et al. \(2018\)](#)). Essentially, all of these papers represent each underlying f_k as a linear combination of a latent set of independent univariate basis functions, which are assigned GP priors. In this note, we revisit the semi-parametric latent factor model (SLFM) of [Teh et al. \(2005\)](#) but construct our latent basis using additive regression trees in the style of [Chipman et al. \(2010\)](#) instead of realizations of GPs.

2 Proposed Model

2.1 Brief Review of BART

[Chipman et al. \(2010\)](#) consider the single task regression problem, modeling $y_i = f(\mathbf{x}_i) + \sigma\epsilon_i$, $\epsilon_i \sim N(0, 1)$ and approximate the unknown function f as a sum of regression trees. To set our notation, let T denote a binary decision tree partitioning \mathbb{R}^p that consists of a collection of interior nodes and $L(T)$ terminal or *leaf* nodes. We associate an axis-aligned decision rule of the form $\{x_j < c\}$ or $\{x_j \geq c\}$ to each internal (i.e. non-leaf) node of T . T defines a partition of \mathbb{R}^p into $L(T)$ rectangular cells and we let $\ell(\mathbf{x}, T)$ be the function that returns the index of the cell containing the point \mathbf{x} . A *regression* tree is a pair (T, M) consisting of

a decision tree T and collection $M = \{\mu_1, \dots, \mu_{L(T)}\}$ of parameters corresponding to each leaf of T . We define the evaluation function $g(\mathbf{x}; T, M) = \mu_{\ell(\mathbf{x}, T)}$ which takes as input a point $\mathbf{x} \in \mathbb{R}^d$ and returns the leaf parameters corresponding to the partition cell containing \mathbf{x} .

At the heart of BART is the approximation

$$f(\mathbf{x}) \approx \sum_{t=1}^m g(\mathbf{x}; T_t, M_t)$$

and a prior over regression trees $\Pi(T, M)$. By modeling each (T_t, M_t) as *a priori* independent realizations from $\Pi(T, M)$, BART implicitly induces a prior over the space of functions from $\mathbb{R}^p \rightarrow \mathbb{R}$. The regression tree prior $\Pi(T, M)$ consists of two parts, a prior $\Pi(T)$ over the space of decisions trees T and a conditional prior $\Pi(M|T)$ of leaf parameters given the decision tree topology. Conditional on the decision tree T , the associated leaf parameters are modeled as i.i.d. $N(\mu_\mu, \sigma_\mu^2)$:

$$\Pi(M|T) = \prod_{\ell=1}^{L(T)} N(\mu_\mu, \sigma_\mu^2).$$

The decision tree prior $\Pi(T)$ corresponds to a branching process and can be described in two parts: the probability that a node at depth d is internal and a distribution over the decision rule at each internal node. Specifically, the probability that node at depth d is internal is $\alpha(1+d)^{-\beta}$ and conditional on a node being internal, the splitting rule is picked uniformly from the set of all available splitting rules. Together, these parts induce a prior over the space of functions $f : \mathbb{R}^d \rightarrow \mathbb{R}$ and we will write $f \sim \text{BART}(m, \alpha, \beta, \mu_\mu, \sigma_\mu)$. [Chipman et al. \(2010\)](#) complete their prior specification by placing a scaled inverse- χ^2 prior over the residual variance $\sigma^2 \sim \frac{\nu\lambda}{\chi_\nu^2}$. We will denote the induced prior on f as $f \sim \text{BART}(m, \alpha, \beta, \mu_\mu, \sigma_\mu^2)$.

Key to the success of BART over a wide variety of applied problems has been the existence of useful *default* choices of the associated hyperparameters. [Chipman et al. \(2010\)](#) recommended setting $\alpha = 0.95$ and $\beta = 2$, which essentially regularizes the depth of the decision trees T_1, \dots, T_m in the BART ensemble. Now that the prior marginal distribution of $\mathbf{f}(\mathbf{x})$ is $N(m\mu_\mu, m\sigma_\mu^2)$. Upon centering and scaling the observed outcomes, we would like this prior to assign substantial prior probability to the range of the standardized data. To this end, [Chipman et al. \(2010\)](#) takes $\mu_\mu = 0$ and set

$$\sigma_\mu = \frac{Y_{\max} - Y_{\min}}{2\kappa\sqrt{m}}$$

where Y_{\max} and Y_{\min} are the maximum and minimum values of the standardized responses. With these choices, the extremes of the observed data, Y_{\max} and Y_{\min} are within 2κ marginal prior standard deviations of $f(\mathbf{x})$ [Chipman et al. \(2010\)](#) recommended $\kappa = 2$ as a good default value. All together, the choices $\alpha = 0.95$, $\beta = 2$, $\mu_\mu = 0$ and $\sigma_\mu = 0.5 \times \kappa^{-1} m^{-1/2} (Y_{\max} - Y_{\min})$ regularize the regression trees in the BART ensembles so that they are not too deep and so that no individual tree accounts for too large a portion of the variation in the observed data. Put another way, BART seeks to fit the observed data well using an ensemble of “weak learners.” With these default choices in mind, abusing our notation slightly, we will take $f \sim \text{BART}(m, \sigma_\mu^2)$ to be a shorthand for $f \sim \text{BART}(m, 0.95, 2, 0, \sigma_\mu^2)$.

2.2 Proposed Model

Throughout, we will assume that the observed data from each task has been centered and scaled to have standard deviation one. In order to perform this scaling, we will require that we observe at least 2 realizations from each task [\[skd\]: This is one disadvantage vis-a-vis using GPs.](#)

In order to induce dependence between the f_k ’s, we follow the basic idea of [?](#) and introduce independent latent basis functions u_1, \dots, u_D and express each f_k as a linear combination of these basis elements:

$$f_k(\mathbf{x}) = \sum_{d=1}^D \phi_{k,d} u_d(\mathbf{x}).$$

More compactly, we will write $\mathbf{f}(\mathbf{x}) = \Phi \mathbf{u}(\mathbf{x})$ where $\mathbf{f}(\mathbf{x}) = (f_1(\mathbf{x}), \dots, f_q(\mathbf{x}))$ and $\mathbf{u}(\mathbf{x}) = (u_1(\mathbf{x}), \dots, u_D(\mathbf{x}))$ and $\Phi = (\phi_{k,d}) \in \mathbb{R}^{q \times D}$ records how much each task f_k depends on the basis element u_d . We place independent $\text{BART}(m, \sigma_\mu^2)$ priors on each of the basis elements u_1, \dots, u_D so that each task f_k is now approximated by a weighted sum of regression trees, which are shared across tasks. This sharing of regression trees induces *a priori* correlation between the tasks. To see this, first note that conditional on Φ ,

$$\begin{aligned} \text{Cov}(f_k(\mathbf{x}), f_{k'}(\mathbf{x}) | \Phi) &= \sum_{d=1}^D \phi_{k,d} \phi_{k',d} \text{Var}(u_d(\mathbf{x})) \\ &= m \sigma_\mu^2 \Phi_{k,\cdot} \Phi_{k',\cdot}^\top, \end{aligned}$$

where $\Phi_{k,\cdot}$ is the $1 \times D$ k^{th} row vector of Φ .

It is worth noting that the first line above holds regardless of the prior placed on the basis elements u_1, \dots, u_D , while the second line follows directly from the fact that the u_d 's are iid $\text{BART}(m, \sigma_\mu^2)$. Now so long as the prior on Φ assigns positive probability to the event $\Phi_{k,\cdot}^\top \Phi_{k',\cdot} \neq 0$, $f_k(\mathbf{x})$ and $f_{k'}(\mathbf{x})$ will be marginally correlated as well.

Just like with univariate BART, we aim to set useful default hyperparameter specifications. To this end, we start with $\sigma_\mu = \frac{Y_{\max} - Y_{\min}}{2\kappa\sqrt{mD}}$ where κ is to be specified, where Y_{\max} and Y_{\min} are the maximum and minimum standardized observations across all observations.

We place independent scaled inverse- χ^2 priors on the residual variances $\sigma_1^2, \dots, \sigma_q^2$, with $\sigma_k^2 \sim \text{Inv. Gamma}(\frac{\nu}{2}, \frac{\nu\lambda_k}{2})$. We use [Chipman et al. \(2010\)](#)'s default choice of $\nu = 3$ and pick λ_k so that *a priori* $\mathbb{P}(\sigma_k < \hat{\sigma}_k) = 0.9$ where $\hat{\sigma}_k$ is an initial over-estimate of the residual standard deviation for task k . With $\nu = 3$, this corresponds to a choice of $\lambda_k = \nu^{-1}q_{\nu,0.1}\hat{\sigma}_k^2$ where $q_{\nu,0.1}$ is the 10% quantile of the χ_ν^2 distribution. In what follows, we will always take $\hat{\sigma}_k = 1$, which is a reasonable over-estimated as we have standardized the responses to have standard deviation 1.

It remains to specify the prior on Φ and to set the number of basis elements D and the number of trees within each basis element m . Recall that one of the key feature of univariate BART is the fact that the marginal prior on $f(\mathbf{x})$ gave substantially probability to the range of the observed responses. While this type of calibration does sacrifice Bayesian coherence, it does ensure that the prior on f is not grossly in conflict with the data. To achieve a similar calibration, notice that conditional on Φ ,

$$\mathbf{f}(\mathbf{x})|\Phi \sim N_q(0_q, m\sigma_\mu^2\Phi\Phi^\top).$$

One way to ensure that the marginal prior on f_k covers the range of observed realizations of the k^{th} task is to require

$$\sqrt{m}\sigma_\mu\|\Phi_{k,\cdot}\|_2 \ll Y_{k,\max} - Y_{k,\min}$$

with high prior probability. Recalling our choice of $\sigma_\mu = \frac{Y_{\max} - Y_{\min}}{2\kappa\sqrt{mD}}$, we can achieve this desideratum if

$$\|\Phi_{k,\cdot}\|_2 \leq \sqrt{D} \times \frac{Y_{k,\max} - Y_{k,\min}}{Y_{\max} - Y_{\min}}, \quad (1)$$

with high prior probability.

As a starting point, consider the relatively simply prior $\phi_{k,d} \sim N(0, \sigma_{\phi,k}^2)$. With this choice, we note that $\|\Phi_{k,\cdot}\|_2^2 \sim \sigma_{\phi,k}^2 \chi_D^2$, so one way to achieve (1) *marginally* for each $k = 1, \dots, q$ is

to set

$$\sigma_{\phi,k} = \left(\frac{D}{q_{D,0.9}} \right)^{1/2} \times \frac{Y_{k,\max} - Y_{k,\min}}{Y_{\max} - Y_{\min}}$$

where $q_{D,0.9}$ is the 90% quantile of a χ_D^2 distribution. Note that we can easily use a different quantile.

A slightly stronger requirement may be to satisfy (1) simultaneously for each task, which motivates taking

$$\sigma_{\phi,k} = \left(\frac{D}{q_{D,0.9^{1/q}}} \right)^{1/2} \times \frac{Y_{k,\max} - Y_{k,\min}}{Y_{\max} - Y_{\min}}.$$

It is worth pointing out that as $D \rightarrow \infty$, the ratio $\frac{D}{q_{D,\tilde{\alpha}}} \leftrightarrow 1$ for any $\tilde{\alpha} > 0.5$. We note that larger choices of $\sigma_{\phi,k}^2$ induce *less* regularization on the fit of f_k since it disperses more marginal probability mass of f_k outside the interval $[Y_{k,\min}, Y_{k,\max}]$.

Note that under this simple prior specification, each task f_k depends on every basis function u_d , and therefore each of the tasks are *a priori* correlated. In some situations, this may be overly restrictive or unrealistic. A more flexible option that can adapt to the situation that some of the tasks are uncorrelated is to use a sparsity inducing prior on the elements of Φ . Specifically, we can introduce binary indicators $\gamma_{k,d}$ and model

$$\begin{aligned} \phi_{k,d} | \gamma_{k,d} &\sim \gamma_{k,d} N(0, \sigma_{\phi,k}^2) + (1 - \gamma_{k,d}) \delta_0 \\ \gamma_{k,1}, \dots, \gamma_{k,d} | \theta_k &\stackrel{\text{i.i.d.}}{\sim} \text{Bernoulli}(\theta_k) \quad \text{for } k = 1, \dots, q \\ \theta_1, \dots, \theta_q &\sim \text{Beta}(a, b) \end{aligned}$$

To set an appropriate value of $\sigma_{\phi,k}^2$, we note that $\sigma_{\phi,k}^{-2} \|\Phi_{k,\cdot}\|_2^2$ is now a χ^2 random variable with a *random* number of degrees of freedom equal to the number of non-zero $\phi_{k,d}$. [skd]: I'm still thinking about the details but I'm pretty sure the number of non-zero entries concentrates around its mean and we can probably do something with that.

2.3 Connection to Existing Work

Our model is closely related to (and in fact, inspired by) the semi-parametric latent factor (SFLM) of ?, who introduce independent basis functions u_1, \dots, u_D and loading matrix $\Phi \in \mathbb{R}^{q \times D}$ and model $\mathbf{f} = \Phi \mathbf{u}$. In that work, they place Gaussian process priors on the

basis functions, $u_d \sim \text{GP}(k_d)$ while have placed BART priors on them. In fact, the SLFM is a special case of the liner model of coregionalization (LMC) that is commonly used for fitting multi-output GPs (see, e.g. [Álvarez et al., 2012](#)). At a high-level, the LMC works by introducing several independent collections of basis functions which are independent draws from a common GP prior and then expressing the functions f_1, \dots, f_q . Typically, the

In the latent factor model $\mathbf{f} = \Phi \mathbf{u}$, we immediately compute the cross-covariance

$$\text{Cov}(f_k(\mathbf{x}), f_{k'}(\mathbf{x}') | \Phi) = \sum_{d=1}^D \phi_{k,d} \phi_{k',d} \text{Cov}(u_d(\mathbf{x}), u_d(\mathbf{x}')).$$

When the basis functions are drawn from Gaussian process priors, $\text{Cov}(u_d(\mathbf{x}), u_d(\mathbf{x}'))$ is just an evaluation of the relevant kernel function. In fact, even in our setting, this covariance also has a convenient kernel representation. Specifically, if $u \sim \text{BART}(m, \sigma_\mu^2)$, then

$$\text{Cov}(u(\mathbf{x}), u(\mathbf{x}')) = m\sigma_\mu^2 \mathbb{P}(\mathbf{x} \sim \mathbf{x}') := m\sigma_\mu^2 k_{\text{BART}}(\mathbf{x}, \mathbf{x}')$$

where we write $\mathbf{x} \sim \mathbf{x}'$ iff \mathbf{x} and \mathbf{x}' are assigned to the same partition cell in a randomly drawn decision tree. This connection between BART and kernel methods was first made in [Linero \(2017\)](#), who further stated a heuristic theorem that so long as $m\sigma_\mu^2 \rightarrow \tilde{\sigma}_\mu^2$ when $m \rightarrow \infty$, a realization from a $\text{BART}(m, \sigma_\mu^2)$ prior converge weakly to a realization from a $\text{GP}(\tilde{\sigma}_\mu^2 k_{\text{BART}})$ prior. So in a certain sense, we may view posterior inference with BART as approximating posterior inference with a $\text{GP}(\tilde{\sigma}_\mu^2 k_{\text{BART}})$ prior.

[Linero et al. \(2018\)](#) is closest in spirit to our model. In that work, they model multiple tasks using a common set of regression trees but introduce a different set of leaf parameters for each task. For a given leaf, the q associated parameters are drawn from a multivariate prior, thereby inducing dependence between the modeled tasks. [\[skd\]: I haven't thought about this too much yet but this model might be a really special case of the one we've described above. In any case, I think the correlation between tasks is fixed throughout and is not learned as it is in our proposal.](#)

3 A Gibbs sampler

Given the data \mathbf{y} and our prior specification, we have a posterior

$$\Pi((T_1^{(1)}, M_1^{(1)}), \dots, (T_m^{(d)}, M_m^{(d)}), \Phi, \boldsymbol{\sigma} | \mathbf{y})$$

on all of the unknown parameters in our semi-parametric latent factor model. We sample from this posterior using a slight elaboration on the backfitting procedure of [Chipman et al. \(2010\)](#). Like their procedure, our's is, at a high level, a Gibbs sampler, that iterates between sequentially updating the basis functions u_1, \dots, u_D , the loading matrix Φ , and the residual variances $\boldsymbol{\sigma}^2$, while keeping the other parameters fixed. We describe each of these conditional updates in the next three subsections.

3.1 Updating the basis functions

Keeping Φ and $\boldsymbol{\sigma}$ fixed, we update each regression tree sequentially holding the remaining $mD - 1$ trees fixed. Each update of a regression tree proceeds in two steps: first, we update the decision tree with a Metropolis-Hastings steps and then draw new leaf parameters conditional on this new decision tree. As we detail below, the MH step is facilitated by an easy-to-compute acceptance probability and the leaf parameters are updated in conjugate fashion.

Recall for each $i = 1, \dots, n$, we have the triplet $(\mathbf{x}_i, \mathbf{y}_i, \delta_i)$ with likelihood model

$$p(\mathbf{y}_i | \Phi, \mathbf{x}_i, \delta_i, \mathbf{u}, \boldsymbol{\sigma}^2) = \prod_{k=1}^q \left[(2\pi\sigma_k^2)^{-\frac{\delta_{i,k}}{2}} \exp \left\{ -\frac{\delta_{i,k} \left(y_{i,k} - \sum_{d=1}^D \phi_{k,d} u_d(\mathbf{x}_i) \right)^2}{2\sigma_k^2} \right\} \right],$$

where $\delta_{i,k} = 1$ if we observe $y_{i,k}$ and 0 otherwise.

Suppose we are updating $(T_t^{(d)}, M_t^{(d)})$, the t^{th} tree of the d^{th} basis function. Define

$$r_{i,k} = y_{i,k} - \sum_{d' \neq d} \phi_{k,d'} u_{d'}(\mathbf{x}_i) - \phi_{k,d} \sum_{t' \neq t} g(\mathbf{x}_i; T_{t'}^{(d)}, M_{t'}^{(d)})$$

when $\delta_{i,k} = 1$ and $r_{i,k} = 0$ otherwise. When $\delta_{i,k} = 1$, $r_{i,k}$ is the partial residual based on the fit of the remaining $mD - 1$ trees. Note that when we do not observe y_i , (i.e. $\delta_{i,k} = 0$), the

specific value of $r_{i,k}$ is immaterial and we have, for convenience, set it to zero. Given the remaining $mD - 1$ trees, Φ, σ , and \mathbf{Y} , the collection of partial residuals $\{r_{i,k}\}$ are sufficient for $(T_t^{(d)}, M_t^{(d)})$:

$$\begin{aligned} p(T, M|\mathbf{y}, \Phi, \sigma, \dots) &= \prod_{\ell=1}^{L(T)} \prod_{i \in I(\ell, T)} \prod_{k=1}^q (2\pi\sigma_k^2)^{-\frac{\delta_{i,k}}{2}} \exp \left\{ -\frac{(r_{i,k} - \phi_{k,d}\mu_\ell)^2}{2\sigma_k^2} \right\} \\ &\times \prod_{\ell=1}^L (2\pi\sigma_\mu^2)^{-\frac{1}{2}} \exp \left\{ -\frac{\mu_\ell^2}{2\sigma_\mu^2} \right\} \\ &\times p(T) \end{aligned}$$

We may write this somewhat more concisely as

$$\begin{aligned} p(T, M|\mathbf{y}, \Phi, \sigma, \dots) &= p(T) \times \prod_{\ell=1}^{L(T)} \left[(2\pi\sigma_\mu^2)^{-\frac{1}{2}} \times \prod_{k=1}^q (2\pi\sigma_k^2)^{-\frac{n_{\ell,k}(T)}{2}} \right] \\ &\times \prod_{\ell=1}^{L(T)} \exp \left\{ -\frac{1}{2} \left[\sigma_\mu^{-2}\mu_\ell^2 + \sum_{k=1}^q \sum_{i \in I(\ell, T)} \delta_{i,k}\sigma_k^{-2}(r_{i,k} - \phi_{k,d}\mu_\ell)^2 \right] \right\} \end{aligned}$$

where $n_{\ell,k}$ counts the number of observations of task k that are associated to leaf ℓ (i.e. $n_{\ell,k} = \sum_{i \in I(\ell, T)} \delta_{i,k}$).

Completing the square, we compute

$$\sigma_\mu^{-2}\mu_\ell^2 + \sum_{k=1}^q \sum_{i \in I(\ell, T)} \delta_{i,k}\sigma_k^{-2}(r_{i,k} - \phi_{k,d}\mu_\ell)^2 = V_\ell^{-1}(\mu - M_\ell)^2 + \sum_{i \in I(\ell, T)} \sum_{k=1}^q \delta_{i,k}\sigma_k^{-2}r_{i,k}^2 - V_\ell^{-1}M_\ell^2,$$

where

$$\begin{aligned} V_\ell &= \left(\sigma_\mu^{-2} + \sum_{k=1}^q n_{\ell,k}\phi_{k,d}^2\sigma_k^{-2} \right)^{-1} \\ M_\ell &= V_\ell \times \sum_{i \in I(\ell, T)} \sum_{k=1}^q \delta_{i,k}\phi_{k,d}r_{i,k}\sigma_k^{-2}. \end{aligned}$$

Integrating out the μ_ℓ 's, we have $p(T|\mathbf{y}, \Phi, \boldsymbol{\sigma}, \dots) \propto$

$$p(T) \times \prod_{\ell=1} \left[V_\ell^{\frac{1}{2}} \sigma_\mu^{-1} \times \left(\prod_{k=1}^q (2\pi\sigma_k^2)^{-\frac{n_{\ell,k}}{2}} \right) \exp \left\{ \frac{1}{2} \left[V_\ell^{-1} M_\ell^2 - \sum_{i \in I(\ell, T)} \sum_{k=1}^q \delta_{i,k} \phi_{k,d} r_{i,k}^2 \right] \right\} \right]$$

Many of the terms above can actually be absorbed into a constant that does not depend on T , leaving us with the less cluttered expression

$$p(T|\mathbf{y}, \Phi, \boldsymbol{\sigma}, \dots) \propto p(T) \times \prod_{\ell=1}^{L(T)} V_\ell^{1/2} \sigma_\mu^{-1} \exp \left\{ \frac{V_\ell^{-1} M_\ell^2}{2} \right\}$$

We also note that

$$p(M|T, \mathbf{y}, \Phi, \boldsymbol{\sigma}, \dots) \propto \prod_{\ell=1}^{L(T)} \exp \left\{ -\frac{(\mu_\ell - M_\ell)^2}{2V_\ell} \right\},$$

enabling us to draw the new leaf parameters $\mu_\ell \sim N(M_\ell, V_\ell)$.

Given a decision tree T , we update it with a simplified version of the transition kernel introduced in [Chipman et al. \(1998\)](#). Specifically, we propose a new tree T^* by either growing T at one terminal node or by pruning T by collapsing two adjacent leafs to their common parent. For a grow proposal T^* , let M_0 and V_0 be the posterior mean and variance of the leaf parameter corresponding to the leaf in T that is split to produce two children in T^* . Similarly, let M_l, M_r, V_l and V_r be the posterior means and variances of the leaf parameter corresponding to these two new leaves in T^* . Then the MH ratio is given by

$$\alpha(T^*, T) = \min \left\{ 1, \frac{q(T^*|T)p(T^*)}{q(T|T^*)p(T)} \times \frac{\left(V_l^{\frac{1}{2}} \sigma_\mu^{-1} \exp \left\{ \frac{M_l^2}{2V_l} \right\} \right) \left(V_r^{\frac{1}{2}} \sigma_\mu^{-1} \exp \left\{ \frac{M_r^2}{2V_r} \right\} \right)}{V_0^{\frac{1}{2}} \sigma_\mu^{-1} \exp \left\{ \frac{M_0^2}{2V_0} \right\}} \right\}$$

We can write down a similar expression for the acceptance probability for prune proposals.

3.2 Updating Φ

We now describe the updates of Φ when we have a Gaussian or spike-and-slab prior on $\phi_{k,d}$.

3.2.1 Gaussian Prior

Observe that

$$p(\Phi_{k,\cdot} | \mathbf{u}, \mathbf{y}, \boldsymbol{\sigma}^2) \propto \exp \left\{ -\frac{1}{2\sigma_{\phi,k}^2} \sum_{d=1}^D \phi_{k,d}^2 - \frac{1}{2\sigma_k^2} \sum_{i=1}^n \delta_{i,k} \left(y_{i,k} - \sum_{d=1}^D \phi_{k,d} u_d(\mathbf{x}_i) \right)^2 \right\}$$

Now if we let

$$V_\Phi = \left[\sigma_\phi^{-2} I_D + \sigma_k^{-2} \sum_{i=1}^n \delta_{i,k} \mathbf{u}(\mathbf{x}_i) \mathbf{u}^\top(\mathbf{x}_i) \right]^{-1}$$

$$M_\Phi = V_\Phi \left[\sum_{i=1}^n \delta_{i,k} y_{i,k} \mathbf{u}(\mathbf{x}_i) \right]$$

it is not difficult to see the posterior distribution of $\Phi_{k,\cdot}$ is $N(M_\Phi, V_\Phi)$. While this expression is tidy, computing V_Φ and M_Φ at each iteration of our Gibbs sampler may be slightly tedious. This is because each of σ_k^2 and $\sum_i \delta_{i,k} \mathbf{u}(\mathbf{x}_i) \mathbf{u}^\top(\mathbf{x}_i)$ are updated at each iteration.

Alternatively, we may update each $\phi_{k,d}$ sequentially, conditional on all of the other elements in Φ . To this end, define $b_i = y_{i,k} - \sum_{d' \neq d} \phi_{k,d'} u_{d'}(\mathbf{x}_i)$. We have

$$p(\phi_{k,d} | \mathbf{u}_{-d}, \mathbf{y}, \boldsymbol{\sigma}^2) \propto \exp \left\{ -\frac{\phi_{k,d}^2}{2\sigma_{\phi,k}^2} - \frac{1}{2\sigma_k^2} \sum_{i=1}^n \delta_{i,k} (b_i - \phi_{k,d} u_d(\mathbf{x}_i))^2 \right\}$$

$$\propto \exp \left\{ -\frac{1}{2} \left[\phi_{k,d}^2 \left[\sigma_{\phi,k}^{-2} + \sigma_k^{-2} \sum_{i=1}^n \delta_{i,k} u_d(\mathbf{x}_i)^2 \right] - 2\phi_{k,d} \sigma_k^{-2} \sum_{i=1}^n \delta_{i,k} b_i \right] \right\}$$

We conclude that $\phi_{k,d} | \dots \sim N(m_\phi, v_\phi)$ where

$$v_\phi = \left[\sigma_{\phi,k}^{-2} + \sigma_k^{-2} \sum_{i=1}^n \delta_{i,k} u_d(\mathbf{x}_i)^2 \right]^{-1} \quad m_\phi = v_\phi \times \sigma_k^{-2} \sum_{i=1}^n \delta_{i,k} b_i$$

3.2.2 Spike-and-slab prior

We re-parametrize like in [Titsias and Lázaro-Gredilla \(2011\)](#) so that $\phi_{k,d} = \gamma_{k,d} \tilde{\phi}_{k,d}$ where *a priori* $\tilde{\phi}_{k,d} \sim N(0, \sigma_{\phi,k}^2)$. For each k , let $\gamma_{k,-d}$ be the set of all indicators associated with task k except for $\gamma_{k,d}$ and similarly define $\tilde{\Phi}_{k,-d}$. We update the pairs $(\tilde{\phi}_{k,d}, \gamma_{k,d})$ holding $\mathbf{u}, \boldsymbol{\sigma}$ and

$\gamma_{k,-d}$ and $\tilde{\Phi}_{k,-d}$ fixed. To this end, let $b_i = y_{i,k} - \sum_{d' \neq d} \phi_{k,d'} u_d(\mathbf{x}_i)$ and note

$$p(\tilde{\phi}_{k,d}, \gamma_{k,d}, |\dots) \propto \exp \left\{ -\frac{1}{2\sigma_k^2} \sum_{i=1}^n \delta_{i,k} (b_i - \gamma_{k,d} \tilde{\phi}_{k,d} u_d(\mathbf{x}_i))^2 \right\} \\ \times \exp \left\{ -\frac{\tilde{\phi}_{k,d}^2}{2\sigma_{\phi,k}^2} \right\} \times \theta_k^{\gamma_{k,d}} (1 - \theta_k)^{1-\gamma_{k,d}}$$

We can define

$$v_\phi(\gamma) = \left(\sigma_\phi^{-2} + \gamma \sigma_k^{-2} \sum_{i=1}^n \delta_{i,k} u_d(\mathbf{x}_i)^2 \right)^{-1} \\ m_\phi(\gamma) = \gamma v_\phi(\gamma) \sigma_k^{-2} \sum_{i=1}^n b_i u_d(\mathbf{x}_i) \delta_{i,k}$$

and marginalize out $\phi_{k,d}$ to obtain the conditional mass function of $\gamma_{k,d}$ given all of the other parameters

$$p(\gamma_{k,d} | \dots) \propto \theta_k^{\gamma_{k,d}} (1 - \theta_k)^{1-\gamma_{k,d}} v_\phi(\gamma_{k,d})^{\frac{1}{2}} \exp \left\{ -\frac{m_\phi(\gamma_{k,d})^2}{2v_\phi(\gamma_{k,d})} \right\}$$

So we may draw a new $\gamma_{k,d}$ according to this distribution. Once we do that, we can draw a new $\tilde{\phi}_{k,d} \sim N(m_\phi(\gamma_{k,d}), v_\phi(\gamma_{k,d}))$ [skd]: the expression of $p(\gamma | \dots)$ above is somewhat different than what Titsias and Lázaro-Gredilla (2011) derived, which is a bit concerning... it's possible I don't follow their notation

3.3 Updating σ

Once we have updated \mathbf{u} and Φ , we can update the residual variances for each task σ_k^2 . We follow Chipman et al. (2010) and model *a priori* $\sigma_k^2 \sim \lambda_k \nu / \chi_\nu^2$. In other words, we have $\sigma_k^2 \sim \text{Inv. Gamma}(\frac{\nu}{2}, \frac{\lambda_k \nu}{2})$. Just as with the rows of Φ , because our likelihood model factorizes over the tasks and because we have independent priors on the task-specific residual variances, we can update the σ_k^2 's independently. Observe that

$$p(\sigma_k^2 | \mathbf{y}, \mathbf{u}, \Phi) \propto (\sigma_k^2)^{-\frac{\nu}{2}-1} \exp \left\{ -\frac{\nu \lambda_k}{2\sigma_k^2} \right\} \times \prod_{i=1}^n (\sigma_k^2)^{-\frac{\delta_{i,k}}{2}} \exp \left\{ -\frac{\delta_{i,k} (y_{i,k} - \Phi \mathbf{u}(\mathbf{x}_i))^2}{2\sigma_k^2} \right\}$$

and we immediately have

$$\sigma_k^2 | \mathbf{y}, \mathbf{u}, \Phi \sim \text{Inv. Gamma} \left(\frac{\nu + \sum \delta_{i,k}}{2}, \frac{\nu \lambda_k + \sum \delta_{i,k} (y_{i,k} - \Phi \mathbf{u}(\mathbf{x}_i))^2}{2} \right)$$

References

- Álvarez, M. A., Rosasco, L., and Lawrence, N. D. (2012). Kernels for vector-valued functions: A review. *Foundations and Trends® in Machine Learning*, 4(3):195–266.
- Breiman, L. and Friedman, J. H. (1997). Predicting multivariate responses in multiple linear regression. *Journal of the Royal Statistical Society (Series B)*, 59(1):3 – 54.
- Cheng, L.-F., Darnell, G., Dumitrascu, B., Chivers, C., Draugelis, M. E., Li, K., and Engelhardt, B. E. (2018). Sparse multi-output gaussian processes for medical time series prediction. arXiv:1703.09112v2.
- Chipman, H. A., George, E. I., and McCulloch, R. E. (1998). Bayesian CART model search. *Journal of the American Statistical Association*, 93(443):935 – 948.
- Chipman, H. A., George, E. I., and McCulloch, R. E. (2010). Bart: Bayesian additive regression trees. *The Annals of Applied Statistics*, 4(1):266 – 298.
- Clifton, L., Clifton, D. A., Pimentel, M. A., Watkinson, P. J., and Tarassenko, L. (2012). Gaussian process regression in vital-sign early warning systems. In *Proceedings of the 34th Annual International Conference of the IEEE EMBS*, pages 6161 – 6164.
- Colopy, G. W., Roberts, S. J., and Clifton, D. A. (2018). Bayesian optimizations of personalized models for patient vital-sign monitoring. *IEEE Journal of Biomedical and Health Informatics*, 22(2):301 – 310.
- Durichen, R., Pimentel, M. A., Clifton, L., Schweikard, A., and Clifton, D. A. (2015). Multitask Gaussian processes for multivariate physiological time-series analysis. *IEEE Transactions on Biomedical Engineering*, 62(1):314 – 322.
- Futoma, J., Hariharan, S., Heller, K., Sendak, M., Brajer, N., Clement, M., Bedoya, A., and O’Brien, C. (2017). An improved multi-output gaussian process rnn with real-time validation for early sepsis detection. In *Proceedings of Machine Learning for Healthcare*.
- Ghassemi, M., Pimentel, M. A., Naumann, T., Brennan, T., Clifton, D. A., Szolovits, P., and Feng, M. (2015). A multivariate timeseries approach to severity of illness assessment and forecasting in ICU with sparse, heterogeneous clinical data. In *Proceedings of the 29th AAAI Conference on Artificial Intelligence (AAAI 2015)*.
- Linero, A. R. (2017). A review of tree-based Bayesian methods. *Communications for Statistical Applications and Methods*, 24(6):543 – 559.

- Linero, A. R., Sinha, D., and Lipsitz, S. R. (2018). Semiparametric mixed-scale models using shared bayesian forests. arXiv:1809.08521.
- Teh, Y. W., Seeger, M., and Jordan, M. I. (2005). Semiparametric latent factor models. In *AISTATS 2005 – Proceedings of the 10th International Workshop on Artificial Intelligence and Statistics*, pages 333 – 340.
- Titsias, M. K. and Lázaro-Gredilla, M. (2011). Spike and slab variational inference for multi-task and multiple kernel learning. In Shawe-Taylor, J., Zemel, R. S., Bartlett, P. L., Pereira, F., and Weinberger, K. Q., editors, *Advances in Neural Information Processing Systems 24*, pages 2339–2347. Curran Associates, Inc.